

# MondgoDB

Kayartaya Vinod

# Introduction

- MongoDB is an open-source document database.
  - Hu**mongo**us **DB**
- Falls into the category of NOSQL databases

# Introduction

- Stores the data in un-normalized format
- Data is stored as collections of documents

# Document database

- A record in an RDBMS is equivalent to a document in MongoDB
  - A data structure composed of field and value pairs.
- MongoDB documents are similar to JSON objects.
  - The values of fields may include scalar data, other documents, arrays, and arrays of documents.
  - Internally, MongoDB stores these documents in the binary format, called BSON (Binary JSON)

# Document

```
{
  id      : 7788,
  name    : "Vinod Kumar",
  phones  : [ "9731424784", "9844083934"],
  emails  : [
    {
      type    : "personal",
      address : "kayartaya.vinod@gmail.com"
    },
    {
      type    : "official",
      address : "vinod@knowledgeworksindia.com"
    }
  ]
}
```

# A Java equivalent of the document

```
public class Email{
    private String type;
    private String address;
    // constructors
    // getters/setters
}

public class Person {
    private int id;
    private String name;
    private String[] phones;
    private List<Email> emails = new ArrayList<Email>();
    // constructors
    // getters/setters
}
```

# A Java equivalent of the document

```
Person p1 = new Person();

p1.setId(7788);
p1.setName("Vinod Kumar");
p1.setPhones(new String[]{"9731424784", "9844093934"});

p1.getEmails().add(
    new Email("personal", "kayartaya.vinod@gmail.com"));
p1.getEmails().add(
    new Email("official", "vinod@knowledgeworksindia.com"));
```

# Document Model

- Data in MongoDB has a flexible schema
- Unlike SQL databases, where you must determine and declare a table's schema before inserting data, MongoDB's collections do not enforce document structure
- This flexibility facilitates the mapping of documents to an entity or an object



# Document Model

- Each document can match the data fields of the represented entity, even if the data has substantial variation
- In practice, however, the documents in a collection share a similar structure
- When designing data models, always consider the application usage of the data (i.e. queries, updates, and processing of the data)

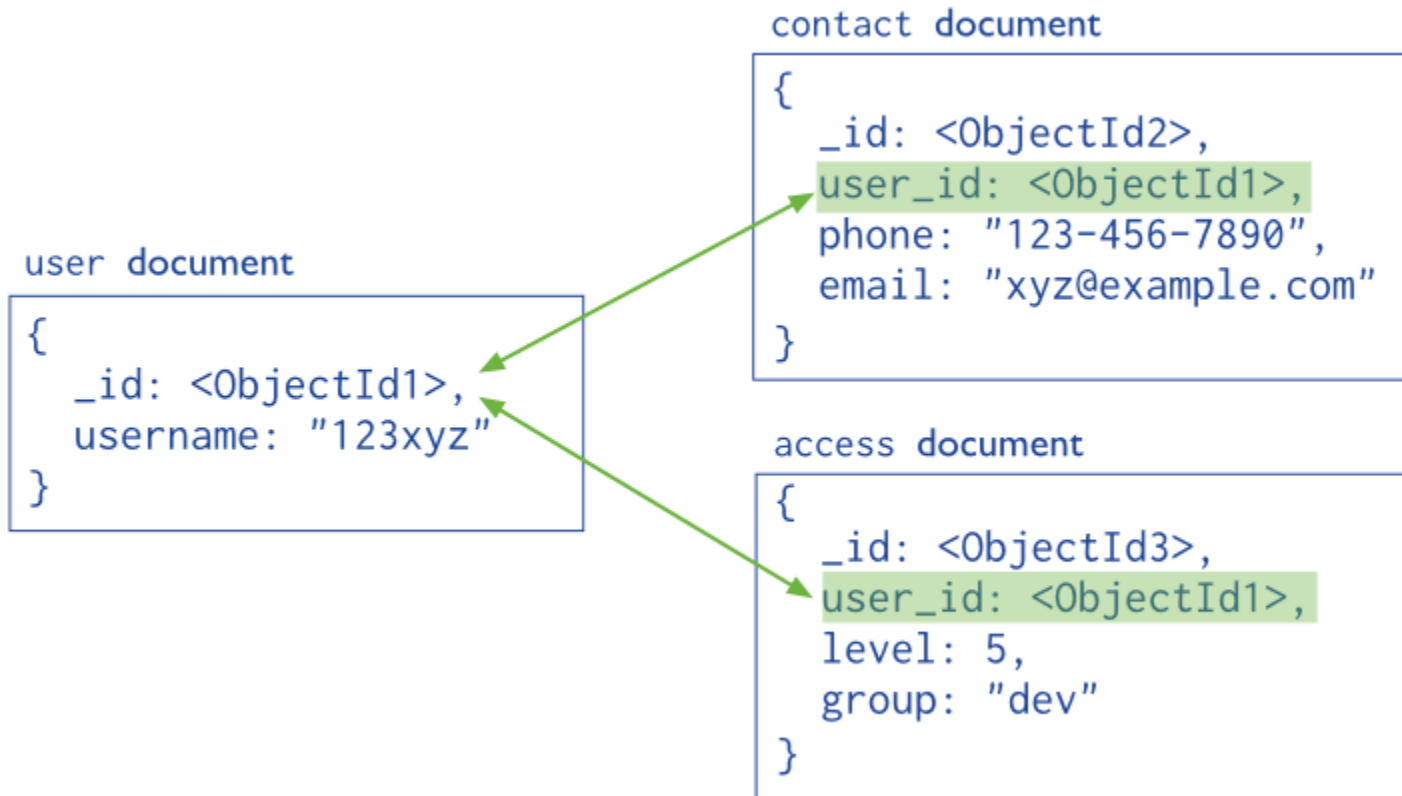
# Document Model

- There are two tools that allow applications to represent these relationships:
  - references and
  - embedded documents

# References

- References store the relationships between data by including links or references from one document to another
- Applications can resolve these references to access the related data
- Broadly, these are normalized data models

# References



# Embedded Data

- Embedded documents capture relationships between data by storing related data in a single document structure
- MongoDB documents make it possible to embed document structures in a field or array within a document
- These denormalized data models allow applications to retrieve and manipulate related data in a single database operation

# References

```
{  
  _id: <ObjectId>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



Embedded sub-document



Embedded sub-document

# Atomicity of Write Operations

- In MongoDB, write operations are atomic at the document level
- No single write operation can atomically affect more than one document or more than one collection
- A denormalized data model with embedded data combines all related data for a represented entity in a single document

# Atomicity of Write Operations

- This facilitates atomic write operations since a single write operation can insert or update the data for an entity
- Normalizing the data would split the data across multiple collections and would require multiple write operations that are not atomic collectively



# Advantages of MongoDB

- Documents correspond to native data types in many programming languages.
- Embedded documents and arrays reduce need for expensive joins.
- Dynamic schema supports fluent polymorphism.

# High performance

- Support for embedded data models reduces I/O activity on database system.
- Indexes support faster queries and can include keys from embedded documents and arrays.

# High availability

- MongoDB's replication facility, called replica sets, provide:
  - automatic failover.
  - data redundancy.
- A replica set is a group of MongoDB servers that maintain the same data set, providing redundancy and increasing data availability.

# Automatic scaling

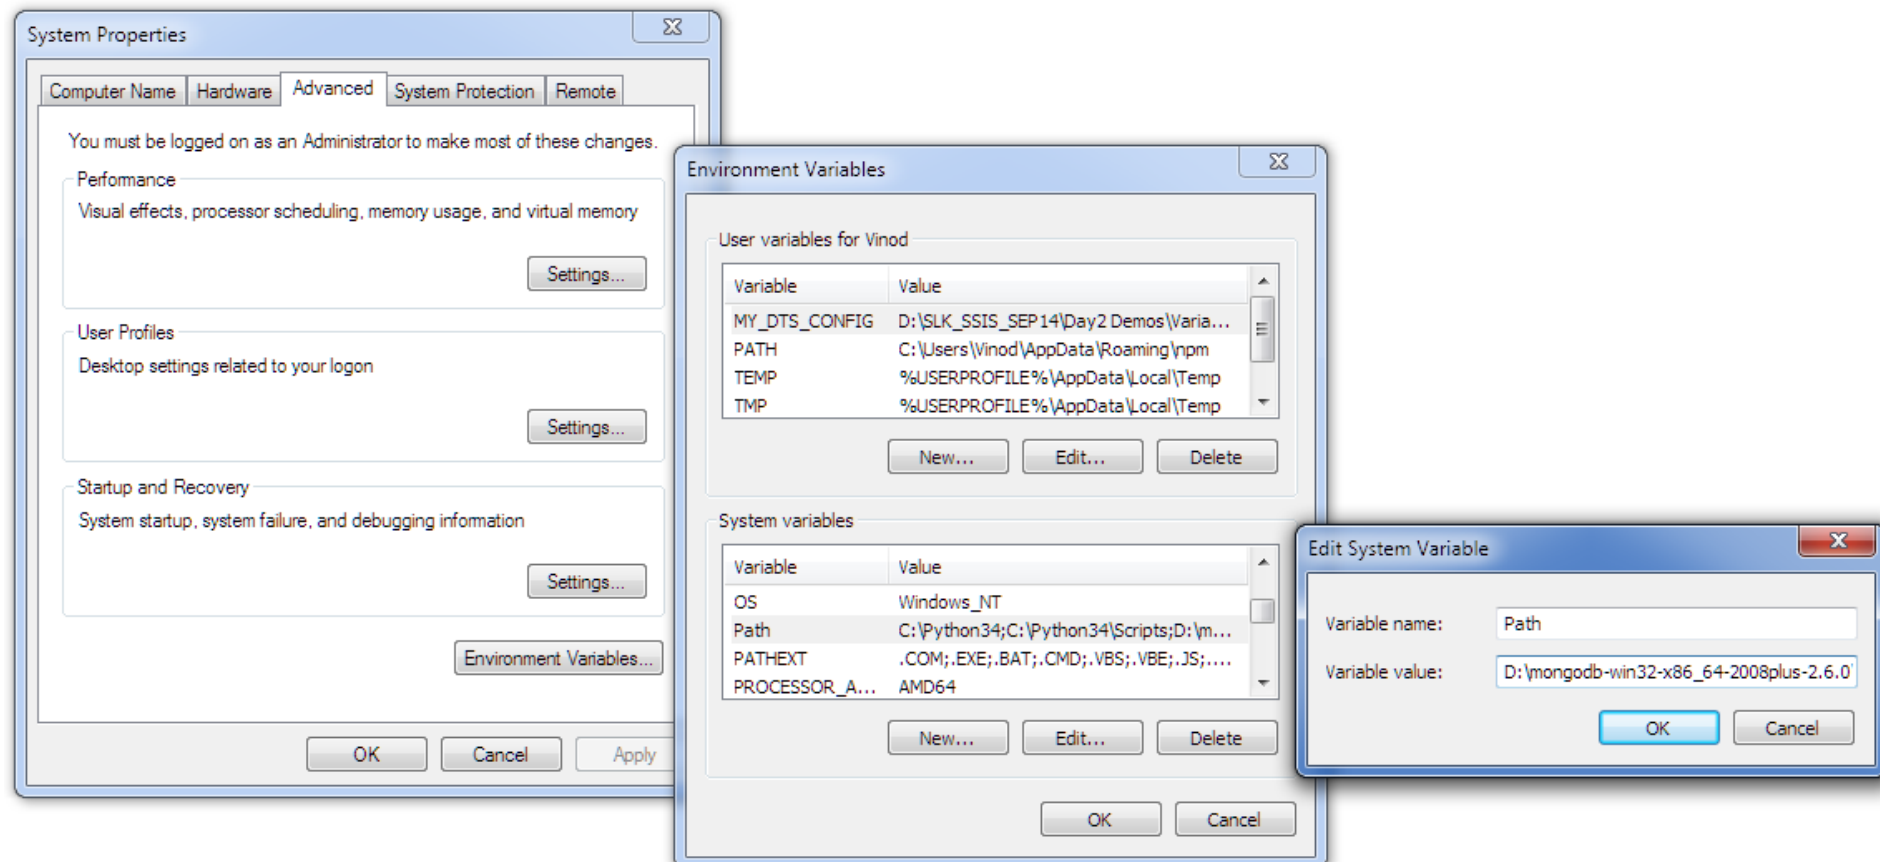
- MongoDB provides horizontal scalability as part of its core functionality.
  - Automatic sharding distributes data across a cluster of machines.
  - Replica sets can provide eventually-consistent reads for low-latency high throughput deployments.

# Installation/setup

- Download the binary for your operating system
  - For Windows 7 64bit:
    - [https://fastdl.mongodb.org/win32/mongodb-win32-x86\\_64-2008plus-2.6.0.zip](https://fastdl.mongodb.org/win32/mongodb-win32-x86_64-2008plus-2.6.0.zip)
  - All Windows downloads:
    - [https://www.mongodb.org/dl/win32/x86\\_64](https://www.mongodb.org/dl/win32/x86_64)
  - All Linux downloads:
    - <https://www.mongodb.org/dl/linux>
- Unzip to a drive
  - In my computer:
    - D:\mongodb-win32-x86\_64-2008plus-2.6.0

# Installation/setup

- Add the D:\mongodb-win32-x86\_64-2008plus-2.6.0\bin to your PATH variable



# Default DB location


- MongoDB requires a data directory to store all data.
- MongoDB's default data directory path is `\data\db`.
  - You can create this folder structure in the same drive as mongod's installation drive
  - Or specify another location when starting the server

# Starting the server

- The server can be started by running the mongod.exe executable
- Some of the useful options are:
  - dbpath=PATH-TO-YOUR-DB
  - port=27017
- Example:  
mongod --dbpath=d:\mongodb-data

(Note: you must create the folder manually before running this command)





```
D:\mongodb-win32-x86_64-2008plus-2.6.0\bin\mongod.exe
2014-04-20T12:27:46.752+0530 [initandlisten] MongoDB starting : pid=1660 port=27017 dbpath=d:\mongodb-data 64-bit host=V
INOD-LENOVO
2014-04-20T12:27:46.753+0530 [initandlisten] targetMinOS: Windows 7/Windows Server 2008 R2
2014-04-20T12:27:46.753+0530 [initandlisten] db version v2.6.0
2014-04-20T12:27:46.753+0530 [initandlisten] git version: 1c1c76aeca21c5983dc178920f5052c298db616c
2014-04-20T12:27:46.753+0530 [initandlisten] build info: windows sys.getwindowsversion(major=6, minor=1, build=7601, pla
tform=2, service_pack='Service Pack 1') BOOST_LIB_VERSION=1_49
2014-04-20T12:27:46.754+0530 [initandlisten] allocator: system
2014-04-20T12:27:46.754+0530 [initandlisten] options: { storage: { dbPath: "d:\mongodb-data" } }
2014-04-20T12:27:46.788+0530 [initandlisten] journal dir=d:\mongodb-data\journal
2014-04-20T12:27:46.789+0530 [initandlisten] recover : no journal files present, no recovery needed
2014-04-20T12:27:46.822+0530 [FileAllocator] allocating new datafile d:\mongodb-data\local.ns, filling with zeroes...
2014-04-20T12:27:46.823+0530 [FileAllocator] creating directory d:\mongodb-data\_tmp
2014-04-20T12:27:46.914+0530 [FileAllocator] done allocating datafile d:\mongodb-data\local.ns, size: 16MB, took 0.09 s
ecs
2014-04-20T12:27:46.917+0530 [FileAllocator] allocating new datafile d:\mongodb-data\local.0, filling with zeroes...
2014-04-20T12:27:47.131+0530 [FileAllocator] done allocating datafile d:\mongodb-data\local.0, size: 64MB, took 0.213 s
ecs
2014-04-20T12:27:47.132+0530 [initandlisten] build index on: local.startup_log properties: { v: 1, key: { _id: 1 }, name
: "_id_", ns: "local.startup_log" }
2014-04-20T12:27:47.133+0530 [initandlisten] added index to empty collection
2014-04-20T12:27:47.160+0530 [initandlisten] command local.$cmd command: create { create: "startup_log", size: 10485760,
capped: true } ntoreturn:1 keyUpdates:0 numYields:0 reslen:37 312ms
2014-04-20T12:27:47.161+0530 [initandlisten] waiting for connections on port 27017
```

# JavaScript Shell

- The executable “mongo.exe” provides an interface to issue direct commands on the db.
- By default the “mongo” command tries to connect to “localhost” and port “27017”
- You can connect to different ones using --host and --port options:

```
mongo --port 12345 --host vinod_homepc
```

D:\mongodb-win32-x86\_64-2008plus-2.6.0\bin\mongo.exe

MongoDB shell version: 2.6.0

connecting to: test

> show dbs

admin (empty)

local 0.078GB

mydb 0.078GB

> use mydb

switched to db mydb

> show collections

persons

system.indexes

> db.persons.findOne()

```
{
  "_id" : ObjectId("535370d8794187ae1c130ee3"),
  "id" : 7788,
  "name" : "Vinod Kumar",
  "phones" : [
    "9731424784",
    "9844083934"
  ],
  "emails" : [
    {
      "type" : "personal",
      "address" : "kayartaya.vinod@gmail.com"
    },
    {
      "type" : "official",
      "address" : "vinod@knowledgeworksindia.com"
    }
  ]
}
```

Run



Type the name of a program, folder, document, or Internet resource, and Windows will open it for you.

Open:

mongo

OK

Cancel

Browse...

# Some commands to start with..

- `show dbs`
  - displays the list of databases
- `use mydb`
  - switches to the database “mydb” if exists, or creates a new with the same name and switches to it
- `db`
  - displays the current database in use
- `db.dropDatabase()`
  - Deletes the current database

# Importing external data

- Use the mongoimport.exe tool to import external data into a database

```
mongoimport
  --host localhost
  --port 27017
  --db mydb
  --jsonArray
  --collection orders
  --file d:\orders.json
```

# Some commands to start with..

- `db.<collection>.findOne()`
  - Displays the first document in the collection
- `db.<collection>.find().pretty()`
  - displays the first 20 documents in an indented format

# Some commands to start with..

- show collections
  - Displays the list of collections (tables in RDBMS)
- db.<collection>.insert(data)
  - If the collection exists, inserts the data, else creates a new collection with the same name and inserts the data

```
p1 = {  
    id      : 6789,  
    name    : "John Doe",  
    city    : "Dallas"  
}  
  
db.persons.insert(p1);
```

```
D:\mongodb-win32-x86_64-2008plus-2.6.0\bin\mongo.exe  
> db.persons.find().pretty()  
{  
  "_id" : ObjectId("535370d8794187ae1c130ee3"),  
  "id" : 7788,  
  "name" : "Vinod Kumar",  
  "phones" : [  
    "9731424784",  
    "9844083934"  
  ],  
  "emails" : [  
    {  
      "type" : "personal",  
      "address" : "kayartaya.vinod@gmail.com"  
    },  
    {  
      "type" : "official",  
      "address" : "vinod@knowledgeworksindia.com"  
    }  
  ]  
}  
{  
  "_id" : ObjectId("53537fbb564dc2b1f4e33401"),  
  "id" : 6789,  
  "name" : "John Doe",  
  "city" : "Dallas"  
}
```



D:\mongodb-win32-x86\_64-2008plus-2.6.0\bin\mongo.exe

```
> show collections
```

```
persons
```

```
system.indexes
```

```
> db.test_data.insert(p1)
```

```
WriteResult({ "nInserted" : 1 })
```

```
> show collections
```

```
persons
```

```
system.indexes
```

```
test_data
```

```
> db.test_data.find().pretty()
```

```
{
```

```
  "_id" : ObjectId("53538053564dc2b1f4e33402"),
```

```
  "id" : 6789,
```

```
  "name" : "John Doe",
```

```
  "city" : "Dallas"
```

```
}
```

```
>
```

# Some commands to start with..

- `db.<collection>.find()`
  - Returns a cursor to the result
  - Displays the first 20 documents on the screen
  - type “it” to iterate again and get 20 more documents

```
> db.salesdata.find()
{ "_id" : 678, "date" : "2014-03-03", "area" : "Jayanagar", "sales" : 11979 }
{ "_id" : 679, "date" : "2014-03-03", "area" : "Basavanagudi", "sales" : 40675 }
{ "_id" : 680, "date" : "2014-03-03", "area" : "Malleshwaram", "sales" : 32669 }
{ "_id" : 681, "date" : "2014-03-03", "area" : "Rajajinagar", "sales" : 32017 }
{ "_id" : 682, "date" : "2014-03-04", "area" : "Jayanagar", "sales" : 11660 }
{ "_id" : 683, "date" : "2014-03-04", "area" : "Basavanagudi", "sales" : 12141 }
{ "_id" : 684, "date" : "2014-03-04", "area" : "Malleshwaram", "sales" : 29496 }
{ "_id" : 685, "date" : "2014-03-04", "area" : "Rajajinagar", "sales" : 16028 }
{ "_id" : 686, "date" : "2014-03-05", "area" : "Jayanagar", "sales" : 23684 }
{ "_id" : 687, "date" : "2014-03-05", "area" : "Basavanagudi", "sales" : 17454 }
{ "_id" : 688, "date" : "2014-03-05", "area" : "Malleshwaram", "sales" : 31525 }
{ "_id" : 689, "date" : "2014-03-05", "area" : "Rajajinagar", "sales" : 19682 }
{ "_id" : 690, "date" : "2014-03-06", "area" : "Jayanagar", "sales" : 26323 }
{ "_id" : 691, "date" : "2014-03-06", "area" : "Basavanagudi", "sales" : 48521 }
{ "_id" : 692, "date" : "2014-03-06", "area" : "Malleshwaram", "sales" : 16901 }
{ "_id" : 693, "date" : "2014-03-06", "area" : "Rajajinagar", "sales" : 37465 }
{ "_id" : 694, "date" : "2014-03-07", "area" : "Jayanagar", "sales" : 12764 }
{ "_id" : 695, "date" : "2014-03-07", "area" : "Basavanagudi", "sales" : 37370 }
{ "_id" : 696, "date" : "2014-03-07", "area" : "Malleshwaram", "sales" : 31562 }
{ "_id" : 697, "date" : "2014-03-07", "area" : "Rajajinagar", "sales" : 12805 }
Type "it" for more
>
```

# Some commands to start with..

- Since “mongo.exe” is a JavaScript shell, you can use a script to process the cursor returned by the find() method

```
var d = db.salesdata.find(); // returns the cursor

// now loop through the cursor to get one document at a time
while(d.hasNext()){
    var s = d.next();
    print(s.date + " >> " + s.area + " Rs." + s.sales);
}
```

D:\mongodb-win32-x86\_64-2008plus-2.6.0\bin\mongo.exe

```
> var d = db.salesdata.find();
> while(d.hasNext()){
... var s = d.next();
... print(s.date + " >> " + s.area + " Rs." + s.sales);
... }
2014-03-03 >> Jayanagar Rs.11979
2014-03-03 >> Basavanagudi Rs.40675
2014-03-03 >> Malleshwaram Rs.32669
2014-03-03 >> Rajajinagar Rs.32017
2014-03-04 >> Jayanagar Rs.11660
2014-03-04 >> Basavanagudi Rs.12141
2014-03-04 >> Malleshwaram Rs.29496
2014-03-04 >> Rajajinagar Rs.16028
2014-03-05 >> Jayanagar Rs.23684
2014-03-05 >> Basavanagudi Rs.17454
2014-03-05 >> Malleshwaram Rs.31525
2014-03-05 >> Rajajinagar Rs.19682
2014-03-06 >> Jayanagar Rs.26323
2014-03-06 >> Basavanagudi Rs.48521
2014-03-06 >> Malleshwaram Rs.16901
2014-03-06 >> Rajajinagar Rs.37465
2014-03-07 >> Jayanagar Rs.12764
2014-03-07 >> Basavanagudi Rs.37370
2014-03-07 >> Malleshwaram Rs.31562
2014-03-07 >> Rajajinagar Rs.12805
```

# Some commands to start with..

- You can use array operator on a cursor returned by find() method

```
var cur = db.salesdata.find();  
var sales1 = cur[3]; // 4th element  
var arr = cur.toArray(); // loads all data to RAM
```

# Performing insert/update

- A new document can be added to a collection using the following methods:
  - `db.<collection>.insert(doc)`
  - `db.<collection>.update(doc, {upsert: true})`
  - `db.<collection>.save(doc)`
- MongoDB does not support transactions
  - Once data is inserted/modified/deleted, it is reflected to all the clients
    - No concept of commit, rollback or savepoints
  - Client applications (such as Java apps) can make use of external transaction managers.

# Performing insert/update

```
D:\mongodb-win32-x86_64-2008plus-2.6.0\bin\mongo.exe
> db.laptops.insert(
... {
...   make: "Lenovo",
...   slno: "CBQ4230641",
...   model: "Z560"
... }
... );
WriteResult({ "nInserted" : 1 })
> db.laptops.find().pretty()
{
  "_id" : ObjectId("5353952bd5a75bbcac27c8eb"),
  "make" : "Lenovo",
  "slno" : "CBQ4230641",
  "model" : "Z560"
}
>
```



# Performing insert/update

- If you add a new document without the `_id` field, the client library or the mongod instance adds an `_id` field and populates the field with a unique ObjectId.
  - A special 12-byte BSON type that guarantees uniqueness within the collection.
  - The ObjectId is generated based on timestamp, machine ID, process ID, and a process-local incremental counter.
  - The `_id` field is immutable

# Performing insert/update

- An existing document can be modified using the update() or save() methods

```
db.<collection>.update(  
  <query>,  
  <update>,  
  {  
    upsert: <boolean>,  
    multi: <boolean>  
  }  
)
```

# Update explained

Parameter	Type	Description
query	document	The selection criteria for the update.
update	document	The modifications to apply.
upsert	boolean	Optional. If set to <code>true</code> , creates a new document when no document matches the query criteria. The default value is <code>false</code> , which does <i>not</i> insert a new document when no match is found.
multi	boolean	Optional. If set to <code>true</code> , updates multiple documents that meet the query criteria. If set to <code>false</code> , updates one document. The default value is <code>false</code> .

# Performing insert/update

```
D:\mongodb-win32-x86_64-2008plus-2.6.0\bin\mongo.exe
> db.laptops.find().pretty()
{
  "_id" : ObjectId("5353952bd5a75bbcac27c8eb"),
  "make" : "Lenovo",
  "slno" : "CBQ4230641",
  "model" : "Z560"
}
> db.laptops.update(
... {make : {$eq: "Lenovo"}},
... {$set :{ model: "Z-560", price: 45000.0}},
... {multi: false}
... )
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.laptops.find().pretty()
{
  "_id" : ObjectId("5353952bd5a75bbcac27c8eb"),
  "make" : "Lenovo",
  "slno" : "CBQ4230641",
  "model" : "Z-560",
  "price" : 45000
}
```

# Update operators

Name	Description
\$inc	Increments the value of the field by the specified amount.
\$mul	Multiplies the value of the field by the specified amount.
\$rename	Renames a field.
\$setOnInsert	Sets the value of a field upon document creation during an upsert. Has no effect on update operations that modify existing documents.
\$set	Sets the value of a field in an existing document.
\$unset	Removes the specified field from an existing document.
\$min	Only updates if the existing field value is less than the specified value.
\$max	Only updates if the existing field value is greater than the specified value.
\$currentDate	Sets the value of a field to current date, either as a Date or a Timestamp.

# Update examples

```
db.laptops.update(  
  {make: {$eq: "Lenovo"}},  
  {  
    $currentDate:{ dop: true },  
    $set: { price: 46500 }  
  }  
)
```

- Adds a new property “dop” with the current date/time as the value
- Changes the value of the property “price” to 46500

# Update examples

```
db.laptops.update(  
  {make: {$eq: "Lenovo"}},  
  {$inc: { price: 1500 }}  
)
```

- Increments the “price” by 1500 for the first matched document

```
db.laptops.update(  
  {make: {$eq: "Lenovo"}},  
  {$inc: { price: 1500 }},  
  {multi: true}  
)
```

- Increments the “price” by 1500 for all the matched documents

D:\mongodb-win32-x86\_64-2008plus-2.6.0\bin\mongo.exe

```
> db.players.find()
{ "_id" : 1, "age" : 28, "name" : "Ravi", "height" : 5.8 }
{ "_id" : 2, "age" : 33, "name" : "Ramesh", "height" : 5.4 }
{ "_id" : 3, "age" : 23, "name" : "Harish", "height" : 5.9 }
{ "_id" : 4, "age" : 55, "name" : "Umesh", "height" : 5.9 }
{ "_id" : 5, "age" : 43, "name" : "Nagesh", "height" : 6.2 }
> db.players.update(
... {age: {$gt: 25}},
... {$inc: {age: 1}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.players.find()
{ "_id" : 1, "age" : 29, "name" : "Ravi", "height" : 5.8 }
{ "_id" : 2, "age" : 33, "name" : "Ramesh", "height" : 5.4 }
{ "_id" : 3, "age" : 23, "name" : "Harish", "height" : 5.9 }
{ "_id" : 4, "age" : 55, "name" : "Umesh", "height" : 5.9 }
{ "_id" : 5, "age" : 43, "name" : "Nagesh", "height" : 6.2 }
> db.players.update( {age: {$gt: 25}}, {$inc: {age: 1}}, {multi: true})
WriteResult({ "nMatched" : 4, "nUpserted" : 0, "nModified" : 4 })
> db.players.find()
{ "_id" : 1, "age" : 30, "name" : "Ravi", "height" : 5.8 }
{ "_id" : 2, "age" : 34, "name" : "Ramesh", "height" : 5.4 }
{ "_id" : 3, "age" : 23, "name" : "Harish", "height" : 5.9 }
{ "_id" : 4, "age" : 56, "name" : "Umesh", "height" : 5.9 }
{ "_id" : 5, "age" : 44, "name" : "Nagesh", "height" : 6.2 }
```



# Update examples

```
db.laptops.update(  
  {make: {$eq: "Lenovo"}},  
  {  
    $rename: {dop: "purchaseDate"},  
    $unset: {price: true}  
  }  
)
```

- Renames the property “dop” to “purchaseDate”
- Removes the property “price” from the document

# Update examples

```
D:\mongodb-win32-x86_64-2008plus-2.6.0\bin\mongo.exe
> db.laptops.findOne()
{
  "_id" : ObjectId("5353952bd5a75bbcac27c8eb"),
  "make" : "Lenovo",
  "sln" : "CBQ4230641",
  "model" : "Z-560",
  "price" : 45000,
  "dop" : ISODate("2014-04-20T10:38:30.289Z")
}
> db.laptops.update(
... {make: {$eq: "Lenovo"}},
... {
... $rename: {dop: "purchaseDate"},
... $unset: {price: true}
... }
... )
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.laptops.findOne()
{
  "_id" : ObjectId("5353952bd5a75bbcac27c8eb"),
  "make" : "Lenovo",
  "sln" : "CBQ4230641",
  "model" : "Z-560",
  "purchaseDate" : ISODate("2014-04-20T10:38:30.289Z")
}
```

# Update examples

```
db.laptops.update(  
  {make: {$eq: "Apple"}},  
  {  
    $set : {  
      slno: "ZXE093745",  
      model: "MacBook Pro",  
      price: 76500  
    }  
  }  
)
```

Only updates the document. If the document does not exist, then no new document is created.

```
db.laptops.update(  
  {make: {$eq: "Apple"}},  
  {  
    $set : {  
      slno: "ZXE093745",  
      model: "MacBook Pro",  
      price: 76500  
    }  
  },  
  { upsert: true }  
)
```

Tries to update the document.

Effect of “upsert: true” →  
If the document does not exist, then a new document is created.

# Deleting a document

- To delete a document from a collection, use a delete criteria and issue the following command:

```
db.<collection>.remove(  
    <query>,  
    <justOne>  
)
```

# Deleting a document

```
db.laptops.remove(  
    { make: "Lenovo"},  
    { justOne: true }  
)
```

- Removes the first matched document

```
db.laptops.remove(  
    { make: "Lenovo"}  
)
```

- Removes all the matched documents

# Finding documents

- For query operations, MongoDB provide a `db.collection.find()` method.
- The method accepts both the query criteria and projections and returns a cursor to the matching documents.
- You can optionally modify the query to impose limits, skips, and sort orders.

# MongoDB query operation

```
db.users.find(  
  { age: 18 },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

# MongoDB query operation

- All queries in MongoDB address a single collection
- You can modify the query to impose limits, skips, and sort orders
- The order of documents returned by a query is not defined unless you specify a sort()
- MongoDB update/remove methods use the same query syntax



# Ordering the query results

- Use the `sort()` method to achieve it

```
db.orders.find({"customer.customer_id": "ANTON"})  
    .sort({order_date: 1})
```

# Query operators

Name	Description
<a href="#"><u>\$gt</u></a>	Matches values that are greater than the value specified in the query.
<a href="#"><u>\$gte</u></a>	Matches values that are equal to or greater than the value specified in the query.
<a href="#"><u>\$in</u></a>	Matches any of the values that exist in an array specified in the query.
<a href="#"><u>\$lt</u></a>	Matches values that are less than the value specified in the query.
<a href="#"><u>\$lte</u></a>	Matches values that are less than or equal to the value specified in the query.
<a href="#"><u>\$ne</u></a>	Matches all values that are not equal to the value specified in the query.
<a href="#"><u>\$nin</u></a>	Matches values that <b>do not</b> exist in an array specified to the query.

# More operators

Name	Description
<a href="#"><u>\$or</u></a>	Joins query clauses with a logical OR returns all documents that match the conditions of either clause.
<a href="#"><u>\$and</u></a>	Joins query clauses with a logical AND returns all documents that match the conditions of both clauses.
<a href="#"><u>\$not</u></a>	Inverts the effect of a query expression and returns documents that do <i>not</i> match the query expression.
<a href="#"><u>\$exists</u></a>	Matches documents that have the specified field.
<a href="#"><u>\$type</u></a>	Selects documents if a field is of the specified type.

# Using \$where operator

```
// Get count of orders having  
// more than 5 products
```

```
db.orders.find(  
    $where: "this.products.length > 5"  
}).count();
```

# Using \$exists operator

```
// Get count of orders having  
// more than 5 products
```

```
db.orders.find(  
    "products.5": {$exists: 1}  
}).count();
```

# Projection

```
db.orders.find(  
  {customer.city: "London"},  
  {  
    _id: 0,  
    order_id: 1,  
    order_date: 1,  
    customer.customer_name: 1,  
    employee.name: 1  
  }) ;
```

# Grouping

- Can be done in several ways
  - aggregate function
  - mapReduce function

# Using aggregate function

- Syntax:

```
db.collection.aggregate(  
    {GROUP_OPTIONS, HAVING_OPTIONS})
```



# Examples

```
db.sales.aggregate(  
{  
  $group: {  
    _id: "$category",  
    salesCount: {$sum: 1}  
  }  
});
```

```
{ "_id" : "Seafood", "salesCount" : 45 }  
{ "_id" : "Produce", "salesCount" : 19 }  
{ "_id" : "Grains/Cereals", "salesCount" : 28 }  
{ "_id" : "Condiments", "salesCount" : 39 }  
{ "_id" : "Meat/Poultry", "salesCount" : 23 }  
{ "_id" : "Dairy Products", "salesCount" : 38 }  
{ "_id" : "Confections", "salesCount" : 48 }  
{ "_id" : "Beverages", "salesCount" : 46 }
```

# Examples

```
db.sales.aggregate(  
  { $group: {  
    _id: "$category",  
    salesCount: { $sum: 1 },  
    salesTotal: { $sum: "$sales" }  
  },  
  { $match: { salesCount: { $gte: 40 } } }  
);
```

```
{ "_id" : "Seafood", "salesCount" : 45, "salesTotal" : 65544.18999999999 }  
{ "_id" : "Confections", "salesCount" : 48, "salesTotal" : 80894.110000000002 }  
{ "_id" : "Beverages", "salesCount" : 46, "salesTotal" : 102074.290000000001 }
```

# Examples

```
db.sales.aggregate({
  $group : {
    _id: "$category",
    profit: { $sum: { $multiply: ["$sales", 0.05] } }
  }
});
```

```
{ "_id" : "Seafood", "profit" : 3277.2095 }
{ "_id" : "Produce", "profit" : 2650.99900000000003 }
{ "_id" : "Grains/Cereals", "profit" : 2797.44100000000007 }
{ "_id" : "Condiments", "profit" : 2763.87800000000006 }
{ "_id" : "Meat/Poultry", "profit" : 4066.90300000000001 }
{ "_id" : "Dairy Products", "profit" : 5737.4875 }
{ "_id" : "Confections", "profit" : 4044.70550000000014 }
{ "_id" : "Beverages", "profit" : 5103.71450000000001 }
```

# Examples- Sorting the group result

```
db.salesdata.aggregate(  
  {  
    $group: {  
      _id: "$quarter",  
      salesCount: { $sum: 1},  
      salesTotal: { $sum: "$amount"},  
      salesAvg: { $avg: "$amount"},  
      maxSales: { $max: "$amount"},  
      minSales: { $min: "$amount"}  
    },  
    {  
      $sort: {  
        salesCount: 1,  
        salesTotal: 1  
      }  
    }  
  }  
)
```

# Operators with \$group

- Following are some of the accumulator operators that could be used along with \$group operator:
  - \$avg, \$first, \$last, \$max, \$min, \$push, \$sum

# Using mapReduce

- Syntax:  
db.collection.mapReduce(  
mapFunction,  
reduceFunction,  
options);

# Using mapReduce

- mapFunction
  - a callback function
  - has access to a single document via “this”
  - should emit two properties from “this” or values derived out of “this” properties

# Using mapReduce

- `emit(a, b)`
  - “a” will be used as a key representing an array of “b” values
- Example:
  - `emit(this.category, this.sales)` will create a dictionary with “category” as key and each “category” representing an array of corresponding “sales” values .



# Example

- Consider the following data

{name: "ram", gender: "male"}

{name: "shyam", gender: "male"}

{name: "sita", gender: "female"}

{name: "gita", gender: "female"}

# Example

- `emit(this.gender, this.name)`
- would create a collection like this

```
[  
  {key: "male", values: [ "ram", "shyam" ] },  
  {key: "female", values: [ "sita", "gita" ] }  
]
```

# Example

- The reduceFunction will receive the key and values from each document separately

# Using mapReduce

- reduceFunction
  - callback function
  - called for each of the key generated by the emit function
  - receives two arguments, key and an array of values

# Example

```
var op=db.persons.mapReduce(function(){  
    var g = "m";  
    var t = "Mr.";  
    if(this.gender=="female"){  
        g = "f";  
        t = "Ms.";  
    }  
    emit(g, t + this.name);  
}, function(k, v){  
    return v.join();  
}, {  
    out: { inline: true}  
});
```

# Example

After calling `emit(g, t + this.name)`

key --> values

"m" --> ["Mr.ram", "Mr.shyam"]

"f" --> ["Ms.sita", "Ms.gita"]

After the `reduceFunction` is called,

key --> values

"m" --> "Mr.ram, Mr.shyam"

"f" --> "Ms.sita, Ms.gita"

# Indexes

- Indexes help efficient execution of queries.
  - Without indexes MongoDB must scan every document in a collection to select those documents that match the query statement.
  - These collection scans are inefficient because they require mongod to process a larger volume of data than an index for each operation.

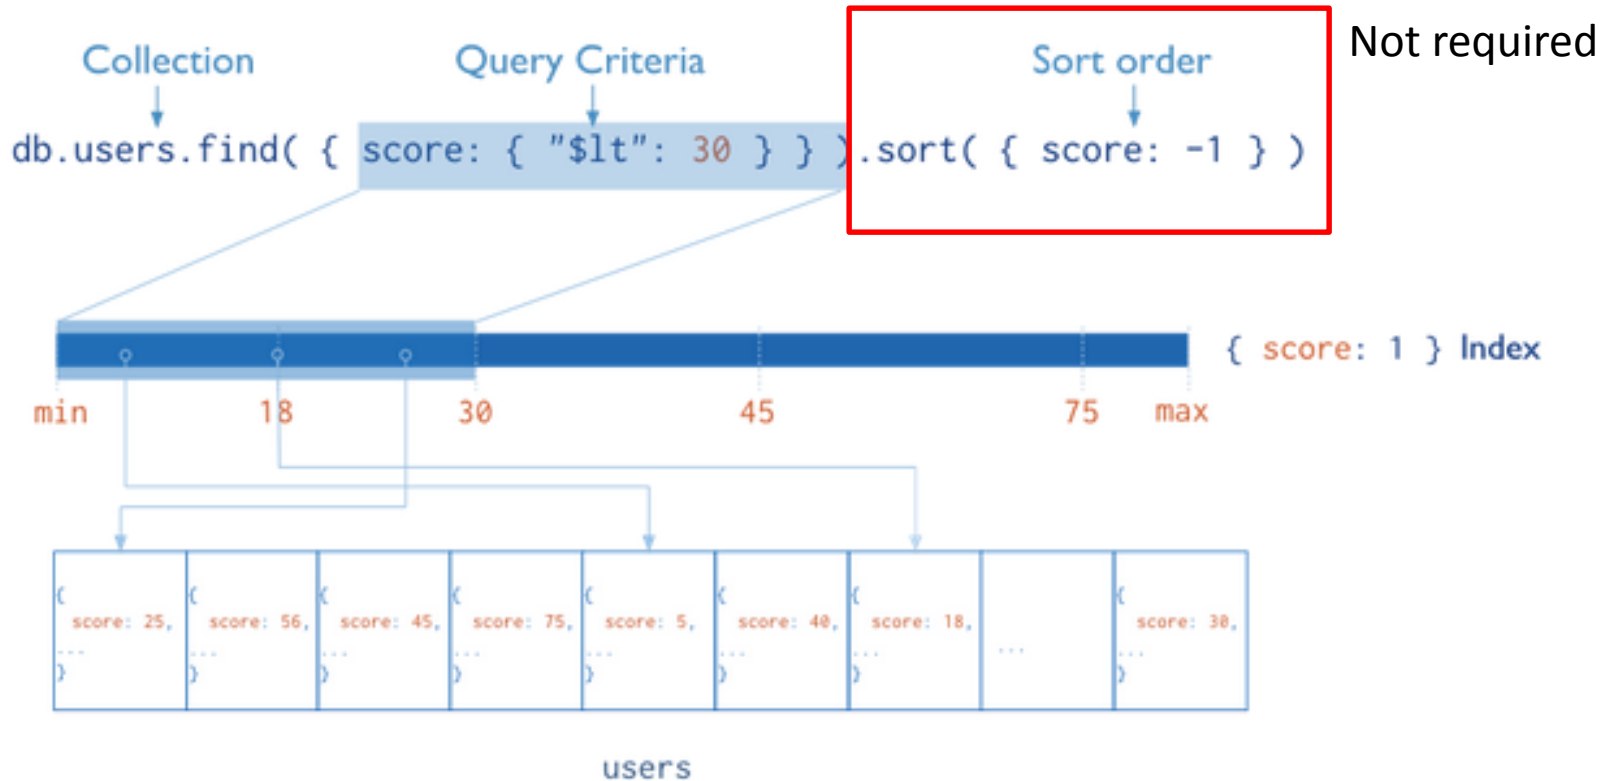
# Indexes

- Indexes are special data structures that store a small portion of the collection's data set in an easy to traverse form.
  - The index stores the value of a specific field or set of fields, ordered by the value of the field.




# Indexes

MongoDB can use indexes to return documents sorted by the index key directly from the index without requiring an additional sort phase.



# Indexes

When the query criteria and the projection of a query include only the indexed fields, MongoDB will return results directly from the index without scanning any documents or bringing documents into memory.



# Index types

- Single field index
- Compound index
- Multi key index
- Geospatial index
- Text index
- Hashed index

<http://docs.mongodb.org/v2.6/core/index-types/>

<http://docs.mongodb.org/manual/core/index-types/>

# Index creation

- Index on single field:

```
db.people.ensureIndex({ "phone-number": 1 })
```

- Compound index

```
db.products.ensureIndex(  
    { item: 1, category: 1, price: 1 })
```

- Text index

```
db.reviews.ensureIndex(  
    { comments: "text" } )
```

# Index creation

- Hashed index

```
db.active.ensureIndex(  
    { productName: "hashed" } )
```

- Geospatial index

```
db.places.ensureIndex({coords: "2dsphere"})  
db.places.ensureIndex({coords: "2d"})
```

# Geospatial Index Example

```
db.atms.ensureIndex({coords: "2dsphere"});
```

```
db.atms.find({  
    coords: {  
        $near: [12.9461, 77.5703]  
    }  
}, {_id: 0}).limit(3).pretty();
```

# Sharding

- Storing data across multiple machines.
- MongoDB uses sharding to support deployments with
  - very large data sets and
  - high throughput operations

# Problems

- Database systems with large data sets and high throughput applications can challenge the capacity of a single server.
  - High query rates can exhaust the CPU capacity of the server.
  - Larger data sets exceed the storage capacity of a single machine.
  - Working set sizes larger than the system's RAM stress the I/O capacity of disk drives.




# Solution 1

- Vertical scaling
  - Increase capacity by adding more CPU and storage resources
    - Limitations: high performance systems with large numbers of CPUs and large amount of RAM are disproportionately more expensive than smaller systems
    - Practical maximum capability
      - Cloud-based providers may only allow users to provision smaller instances


# Solution 2

- Sharding
  - Horizontal scaling
  - Divides the data set and distributes the data over multiple servers, or shards.
    - Each shard is an independent database,
    - Collectively, the shards make up a single logical database.

# Sharding



# Sharding

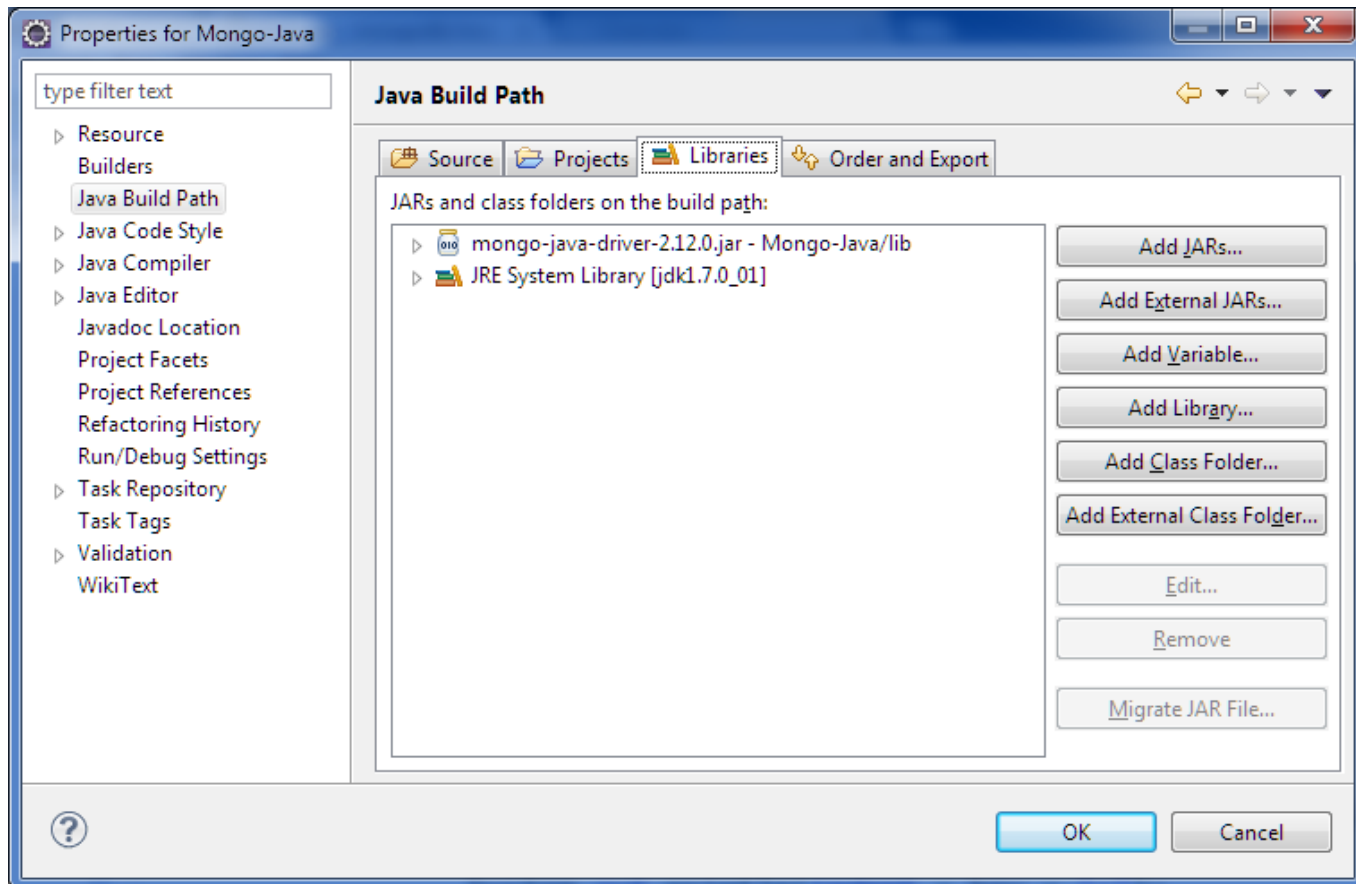


# Using JavaDriver

- Download the latest version of JavaDriver
  - <https://github.com/mongodb/mongo-java-driver/releases/download/r2.12.0/mongo-java-driver-2.12.0.jar>

# Using JavaDriver

- Add to buildpath (classpath)



# Use the API

- `com.mongodb.MongoClient`
  - A MongoDB client with internal connection pooling.
  - For most applications, you should have one `MongoClient` instance for the entire JVM.
  - The `MongoClient` class is designed to be thread safe and shared among threads
  - Inherits from `com.mongodb.Mongo`

# Connecting to server

```
MongoClient client = new MongoClient();
```

```
MongoClient client = new MongoClient("localhost");
```

```
MongoClient client = new MongoClient("localhost", 27017);
```

```
MongoClient client = new MongoClient(  
    new ServerAddress("localhost"));
```

```
MongoClient client = new MongoClient(  
    new ServerAddress("localhost"),  
    new MongoClientOptions.Builder().build());
```



# Connecting to replica set

- You can connect to a replica set using the Java driver by passing a `ServerAddress` list to the `Mongo` constructor

```
Mongo mongo = new Mongo(Arrays.asList(  
    new ServerAddress("localhost", 27017),  
    new ServerAddress("localhost", 27018),  
    new ServerAddress("localhost", 27019))  
);
```

# Some useful methods

- `List<String> getDatabaseNames()`
- `DB getDB(String dbName)`
- `void dropDatabase(String dbName)`
- `void close()`

# Query for list of databases

```
MongoClient client = new MongoClient("localhost", 27017);
List<String> dbNames = client.getDatabaseNames();

System.out.println("Following databases were found: ");
for (String dbName : dbNames) {
    System.out.println(dbName);
}

client.close();
```

# Mongo Database

- `com.mongodb.DB`
- A thread-safe client view of a logical database in a MongoDB cluster.

```
MongoClient client = new MongoClient();  
DB db = client.getDB("mydb");
```

# Some useful methods

- DBCollection createCollection(  
    String name,DBObject options)
- DBCollection getCollection(String name)
- Set<String> getCollectionNames()
- void dropDatabase()

# Query for list of collections

```
MongoClient client = new MongoClient("localhost", 27017);
DB db = client.getDB("mydb");
Set<String> collections = db.getCollectionNames();

System.out.println("'mydb' contains following
collections");
for (String collection : collections) {
    System.out.println(collection);
}
client.close();
```

# Mongo Collection

- `com.mongodb.DBCollection`
- This class provides a skeleton implementation of a database collection

```
MongoClient client = new MongoClient();  
DB db = client.getDB("mydb");  
DBCollection orders = db.getCollection("orders");
```

# Adding a document

```
MongoClient mc = new MongoClient();  
DB db = mc.getDB("vindb");
```

```
DBObject doc1 = new BasicDBObject("name", "Cellphone")  
    .append("price", 25000.0)  
    .append("make", "Lenovo");
```

```
DBCollection items = db.getCollection("items");  
items.save(doc1);
```



# Adding a json string

```
MongoClient mc = new MongoClient();
```

```
DB db = mc.getDB("vindb");
```

```
String itemStr = "{" +  
    "name : \"Wrist Watch\", " +  
    "make : \"Titan\", " +  
    "price: 5600.0" +  
    "}";
```

```
DBObject doc1 = (DBObject) JSON.parse(itemStr);
```

```
DBCollection items = db.getCollection("items");
```

```
items.save(doc1);
```

# Getting data

```
MongoClient mc = new MongoClient();  
DB db = mc.getDB("vindb");  
DBCollection sales = db.getCollection("sales");  
  
DBObject first = sales.findOne();  
  
System.out.println(first);  
  
System.out.println("Sales amount = $" + first.get("sales"));  
System.out.println("Quarter    = " + first.get("quarter"));  
System.out.println("Category   = " + first.get("category"));  
System.out.println("Name      = " + first.get("product"));
```

# Get all data

```
MongoClient mc = new MongoClient();  
DB db = mc.getDB("vindb");  
DBCollection books = db.getCollection("books");  
DBCursor cursor = books.find();  
  
String booksJson = JSON.serialize(cursor);  
System.out.println(booksJson);
```

# Querying

```
MongoClient mc = new MongoClient();  
DB db = mc.getDB("vindb");  
DBCollection books = db.getCollection("sales");  
  
DBObject query = new BasicDBObject(  
    "category", "Beverages");  
  
DBCursor cursor = books.find(query);
```